


```
# import some basic Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
```

<https://archive.ics.uci.edu/dataset/464/superconductivity+data>

```
# loading the dataset
from google.colab import files
file = files.upload()
df = pd.read_csv('train.csv')
df.head()
```

  No file chosen      Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.


Saving train.csv to train (1).csv

	number_of_elements	mean_atomic_mass	wtd_mean_atomic_mass	gmean_atomic_mass	wtd_gmean_atomic_mass	entropy_atomic_mass	wtd_entropy
0	4	88.944468	57.862692	66.361592	36.116612	1.181795	
1	5	92.729214	58.518416	73.132787	36.396602	1.449309	
2	4	88.944468	57.885242	66.361592	36.122509	1.181795	
3	4	88.944468	57.873967	66.361592	36.119560	1.181795	
4	4	88.944468	57.840143	66.361592	36.110716	1.181795	

5 rows x 82 columns

## ✓ Data Cleaning & Preprocessing

```
# Checking basic info and check for missing values
print("Dataset Info:")
df.info()
print("\nMissing Values:")
print(df.isnull().sum())
```

```

15  entropy_fie          21263 non-null  float64
16  wtd_entropy_fie     21263 non-null  float64
17  range_fie           21263 non-null  float64
18  wtd_range_fie       21263 non-null  float64
19  std_fie             21263 non-null  float64
20  wtd_std_fie         21263 non-null  float64
21  mean_atomic_radius  21263 non-null  float64
22  wtd_mean_atomic_radius 21263 non-null  float64
23  gmean_atomic_radius 21263 non-null  float64
24  wtd_gmean_atomic_radius 21263 non-null  float64
25  entropy_atomic_radius 21263 non-null  float64
26  wtd_entropy_atomic_radius 21263 non-null  float64
27  range_atomic_radius 21263 non-null  int64
28  wtd_range_atomic_radius 21263 non-null  float64
```

```

50 wtd_std_FusionHeat          21263 non-null float64
51 mean_FusionHeat             21263 non-null float64
52 wtd_mean_FusionHeat         21263 non-null float64
53 gmean_FusionHeat            21263 non-null float64
54 wtd_gmean_FusionHeat        21263 non-null float64
55 entropy_FusionHeat          21263 non-null float64
56 wtd_entropy_FusionHeat      21263 non-null float64
57 range_FusionHeat            21263 non-null float64
58 wtd_range_FusionHeat        21263 non-null float64
59 std_FusionHeat              21263 non-null float64
60 wtd_std_FusionHeat          21263 non-null float64
61 mean_ThermalConductivity    21263 non-null float64
62 wtd_mean_ThermalConductivity 21263 non-null float64
63 gmean_ThermalConductivity   21263 non-null float64
64 wtd_gmean_ThermalConductivity 21263 non-null float64
65 entropy_ThermalConductivity 21263 non-null float64
66 wtd_entropy_ThermalConductivity 21263 non-null float64
67 range_ThermalConductivity   21263 non-null float64
68 wtd_range_ThermalConductivity 21263 non-null float64
69 std_ThermalConductivity     21263 non-null float64
70 wtd_std_ThermalConductivity 21263 non-null float64
71 mean_Valence                21263 non-null float64
72 wtd_mean_Valence            21263 non-null float64
73 gmean_Valence               21263 non-null float64

```

```

# Summary statistics
print("\nSummary Statistics:")
print(df.describe())

```



#### Summary Statistics:

	number_of_elements	mean_atomic_mass	wtd_mean_atomic_mass \
count	21263.000000	21263.000000	21263.000000
mean	4.115224	87.557631	72.988310
std	1.439295	29.676497	33.490406
min	1.000000	6.941000	6.423452
25%	3.000000	72.458076	52.143839
50%	4.000000	84.922750	60.696571
75%	5.000000	100.404410	86.103540
max	9.000000	208.980400	208.980400

	gmean_atomic_mass	wtd_gmean_atomic_mass	entropy_atomic_mass \
count	21263.000000	21263.000000	21263.000000
mean	71.290627	58.539916	1.165608
std	31.030272	36.651067	0.364930
min	5.320573	1.960849	0.000000
25%	58.041225	35.248990	0.966676
50%	66.361592	39.918385	1.199541
75%	78.116681	73.113234	1.444537
max	208.980400	208.980400	1.983797

	wtd_entropy_atomic_mass	range_atomic_mass	wtd_range_atomic_mass \
count	21263.000000	21263.000000	21263.000000
mean	1.063884	115.601251	33.225218
std	0.401423	54.626887	26.967752
min	0.000000	0.000000	0.000000
25%	0.775363	78.512902	16.824174
50%	1.146783	122.906070	26.636008
75%	1.359418	154.119320	38.356908
max	1.958203	207.972460	205.589910

	std_atomic_mass ...	wtd_mean_Valence	gmean_Valence \
count	21263.000000 ...	21263.000000	21263.000000
mean	44.391893 ...	3.153127	3.056536
std	20.035430 ...	1.191249	1.046257
min	0.000000 ...	1.000000	1.000000
25%	32.890369 ...	2.116732	2.279705
50%	45.123500 ...	2.618182	2.615321
75%	59.322812 ...	4.026201	3.727919
max	101.019700 ...	7.000000	7.000000

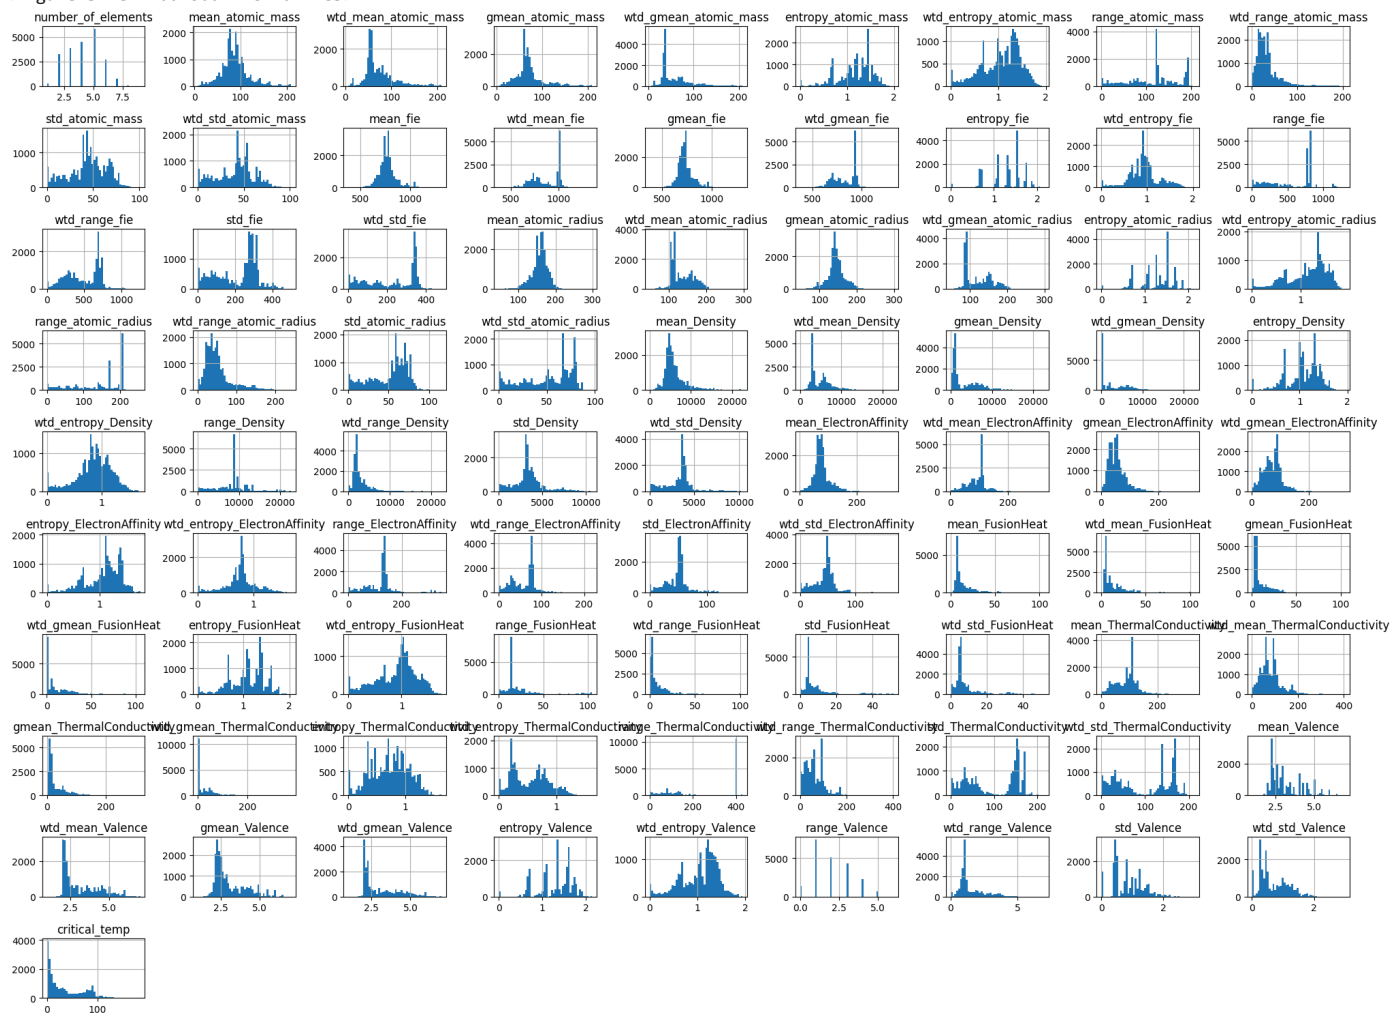
	wtd_gmean_Valence	entropy_Valence	wtd_entropy_Valence	range_Valence \
count	21263.000000	21263.000000	21263.000000	21263.000000
mean	3.055885	1.295682	1.052841	2.041010
std	1.174815	0.393155	0.380291	1.242345
min	1.000000	0.000000	0.000000	0.000000
25%	2.091251	1.060857	0.775678	1.000000
50%	2.434057	1.368922	1.166532	2.000000
75%	3.914868	1.589027	1.330801	3.000000
max	7.000000	2.141963	1.949739	6.000000

	wtd_range_Valence	std_Valence	wtd_std_Valence	critical_temp
count	21263.000000	21263.000000	21263.000000	21263.000000
mean	1.483007	0.839342	0.673987	34.421219
std	0.978176	0.484676	0.455580	34.254362


	min	0.000000	0.000000	0.000000	0.000210
25%	0.921454	0.451754	0.306892	5.365000	

```
# Plot histograms of features
plt.figure(figsize=(12, 8))
df.hist(bins=50, figsize=(20, 15))
plt.tight_layout()
plt.show()
```

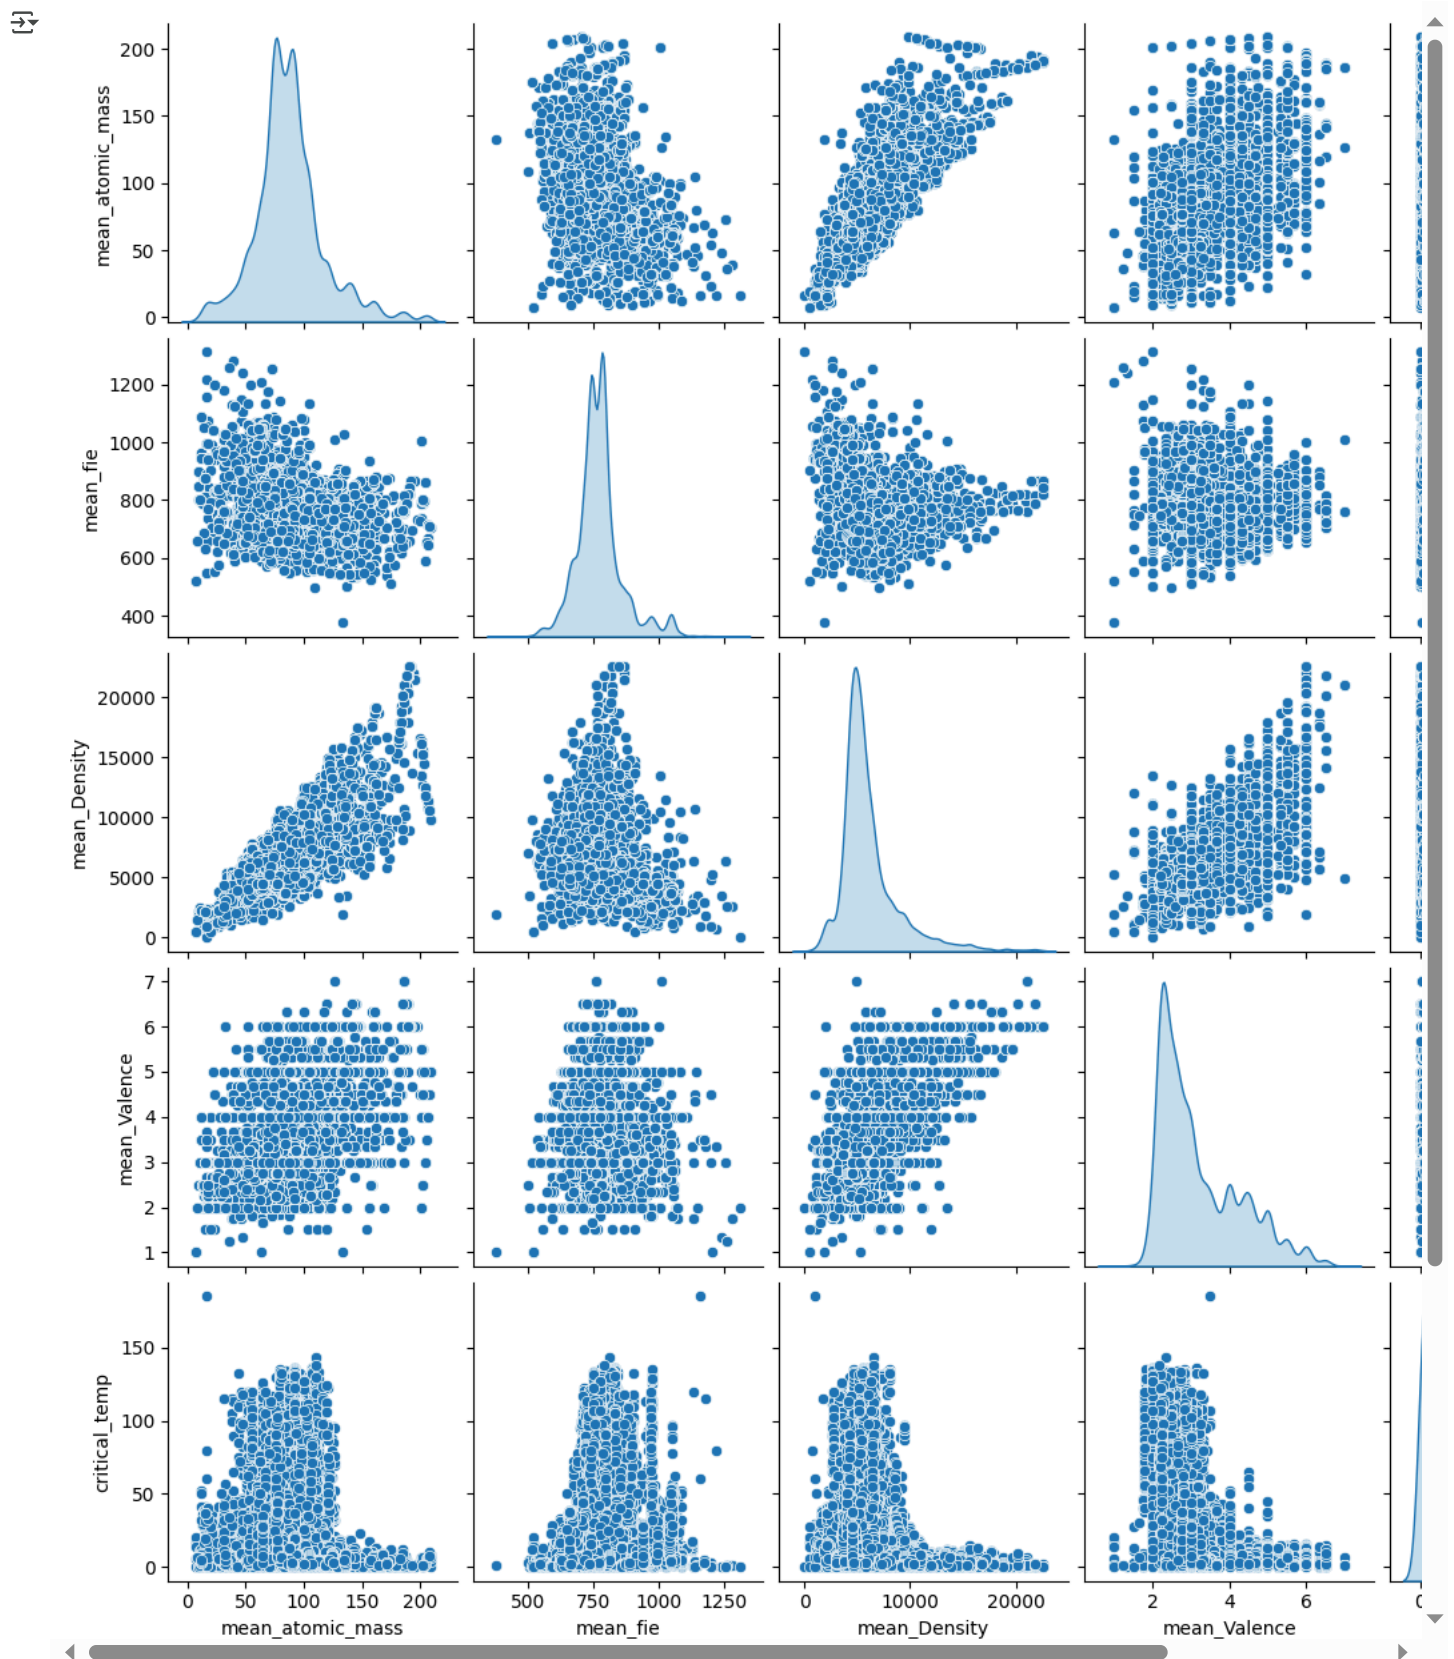
↩ <Figure size 1200x800 with 0 Axes>



```
# Correlation matrix heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), cmap='coolwarm', annot=False, linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
# Pair plot of selected key features
selected_features = ['mean_atomic_mass', 'mean_fie', 'mean_Density', 'mean_Valence', 'critical_temp']
sns.pairplot(df[selected_features], diag_kind='kde')
plt.show()
```



## ✓ Principle Component Analysis

```
# Drop the target variable (critical_temp) since PCA is unsupervised
X = df.drop(columns=['critical_temp'])
```

```
# Standardizing the data
scaler = StandardScaler()
```

```

X_scaled = scaler.fit_transform(X)


# Convert back to DataFrame for better readability
X_scaled_df = pd.DataFrame(X_scaled, columns=X.columns)

# Apply PCA
pca = PCA()
X_pca = pca.fit_transform(X_scaled)

# Explained variance ratio
explained_variance = pca.explained_variance_ratio_

# Plot cumulative explained variance
plt.figure(figsize=(8, 5))
plt.plot(np.cumsum(explained_variance), marker='o', linestyle='--')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('PCA - Explained Variance')
plt.grid()
plt.show()

```



```

# Select number of components (e.g., 95% variance)
num_components = np.argmax(np.cumsum(explained_variance) >= 0.95) + 1
print(f"Optimal number of components: {num_components}")

# Apply PCA with selected components
pca_final = PCA(n_components=num_components)
X_pca_final = pca_final.fit_transform(X_scaled)

# Convert PCA results into a DataFrame
X_pca_df = pd.DataFrame(X_pca_final, columns=[f'PC{i+1}' for i in range(num_components)])

```

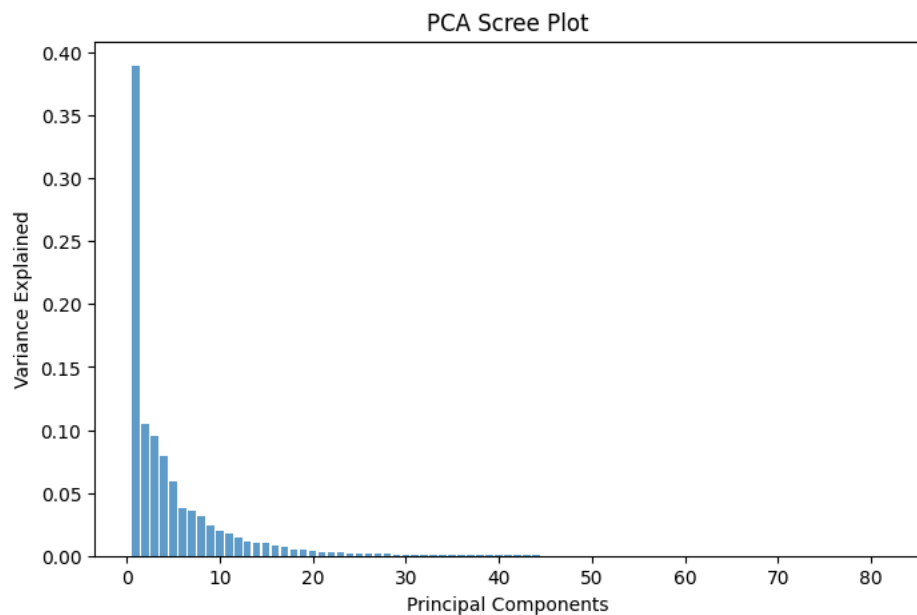


Optimal number of components: 17


```

plt.figure(figsize=(8, 5))
plt.bar(range(1, len(explained_variance) + 1), explained_variance, alpha=0.7)
plt.xlabel('Principal Components')
plt.ylabel('Variance Explained')
plt.title('PCA Scree Plot')
plt.show()

```



```
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_pca_df['PC1'], y=X_pca_df['PC2'])
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA Scatter Plot')
plt.grid()
plt.show()
```



```
# Drop the target variable (critical_temp)
X = df.drop(columns=['critical_temp'])

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

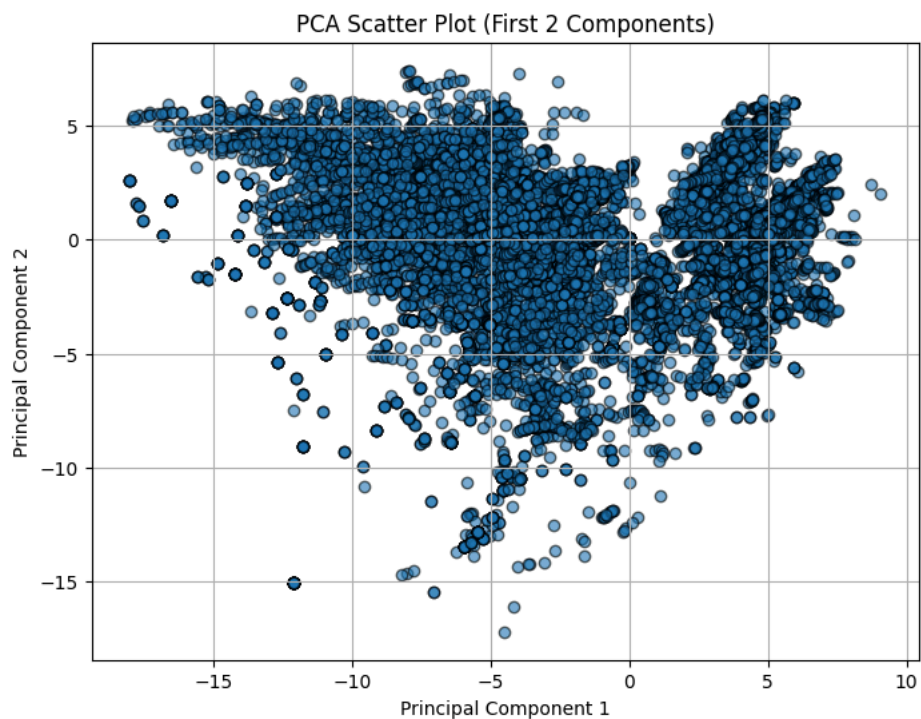

# Apply PCA with optimal 17 components
pca = PCA(n_components=17)
X_pca = pca.fit_transform(X_scaled)
```



```
# Explained variance ratio for 17 components
explained_variance = pca.explained_variance_ratio_
cumulative_variance = np.cumsum(explained_variance)

# Plot cumulative explained variance
plt.figure(figsize=(8, 5))
plt.plot(range(1, 18), cumulative_variance, marker='o', linestyle='--', color='b')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Cumulative Explained Variance for 17 Components')
plt.grid()
plt.show()

# Scatter plot for first two principal components
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], alpha=0.6, edgecolors='k')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA Scatter Plot (First 2 Components)')
plt.grid()
plt.show()
```





## ✓ Linear Regression Model vs Random Forest for Model Performance

```
# Import necessary libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Step 1: Define target variable and features
y = df['critical_temp']
X = df.drop(columns=['critical_temp'])

# Step 2: Apply PCA to features (X)
# (Assuming PCA has already been applied to X, resulting in X_pca)

# Step 3: Split the PCA-transformed data
X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=0.2, random_state=42)

# Step 4: Train your model (e.g., Linear Regression)

# Step 2: Train a Linear Regression Model
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

# Step 3: Make Predictions
y_pred = lin_reg.predict(X_test)

# Step 4: Evaluate Model Performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Linear Regression Model Performance:")
print(f"📊 Mean Squared Error (MSE): {mse:.2f}")
print(f"📈 R² Score: {r2:.4f}")

# Step 5: Try a More Advanced Model - Random Forest
rf_reg = RandomForestRegressor(n_estimators=100, random_state=42)
rf_reg.fit(X_train, y_train)

# Predictions
y_pred_rf = rf_reg.predict(X_test)

# Evaluate
mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print(f"\nRandom Forest Model Performance:")
print(f"📊 Mean Squared Error (MSE): {mse_rf:.2f}")
print(f"📈 R² Score: {r2_rf:.4f}")
```

🔄 Linear Regression Model Performance:  
📊 Mean Squared Error (MSE): 446.48  
📈 R² Score: 0.6121

Random Forest Model Performance:  
📊 Mean Squared Error (MSE): 95.39  
📈 R² Score: 0.9171

### Summary of the Analysis:

- The Random Forest model significantly outperforms the Linear Regression model, with a lower MSE and a much higher R² score.
- This suggests that Random Forest is better at capturing complex patterns and relationships in the data.
- The Linear Regression model, while decent, is less effective in predicting the target variable, as it doesn't capture as much of the variance compared to Random Forest.

```
from sklearn.model_selection import RandomizedSearchCV

# Define the Random Forest model
rf_model = RandomForestRegressor(random_state=42)
```

```
# Define the parameter distribution for randomized search
param_dist = {
    'n_estimators': [100, 200, 300, 400],
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10]
}

# Perform randomized search with n_iter set to a smaller value to speed up
random_search = RandomizedSearchCV(rf_model, param_distributions=param_dist, n_iter=10, cv=3, n_jobs=-1, random_state=42)
random_search.fit(X_train, y_train)

# Get the best model and print the results
print("Best Parameters:", random_search.best_params_)
rf_best = random_search.best_estimator_

# Evaluate the best model
y_pred_rf = rf_best.predict(X_test)
mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print(f"Random Forest Test MSE: {mse_rf}, R²: {r2_rf}")
```

```
➡ Best Parameters: {'n_estimators': 400, 'min_samples_split': 2, 'max_depth': 30}
Random Forest Test MSE: 94.31938522021353, R²: 0.9180601278578859
```

**The results of the RandomizedSearchCV indicate that the best hyperparameters for your Random Forest model are:**

- n\_estimators: 400 (number of trees in the forest)
- min\_samples\_split: 2 (the minimum number of samples required to split an internal node)
- max\_depth: 30 (the maximum depth of the trees)
- Model Performance:
  - Test MSE (Mean Squared Error): 94.32
  - R² (R-squared): 0.9181

#### Interpretation of the Results:

- MSE (94.32): A relatively low MSE suggests that the model's predictions are fairly close to the actual values. Since the MSE is based on squared errors, a lower MSE means better model performance.
- R² (0.9181): This value indicates that the model explains approximately 91.81% of the variance in the critical temperature of superconductors. This is a strong result, as an R² score close to 1 means the model is doing an excellent job of fitting the data.

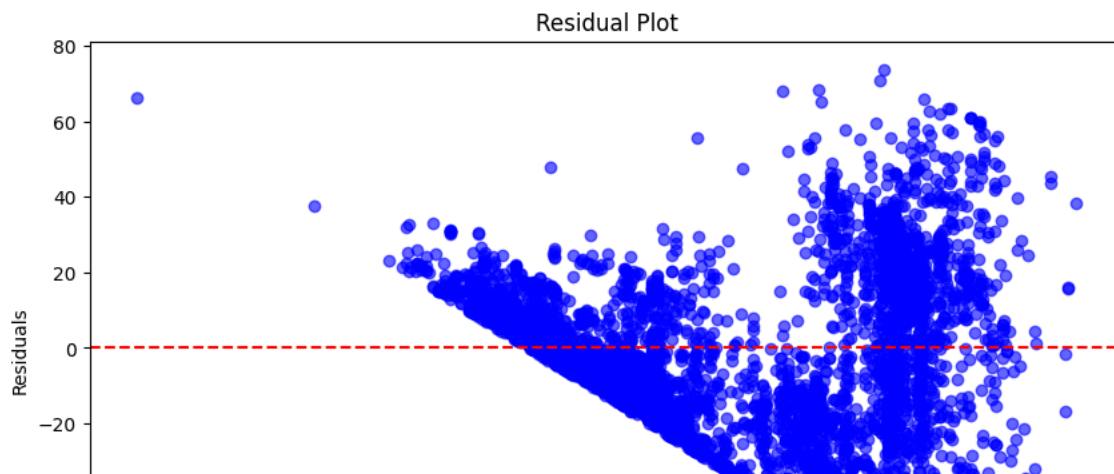
#### Residual Plot

- This plot helps you understand how well your model's predictions match the actual values. Ideally, the residuals should be scattered randomly around zero, which suggests no patterns are left to be captured by the model.

```
import matplotlib.pyplot as plt

# Calculate residuals
residuals = y_test - y_pred

# Plotting Residuals
plt.figure(figsize=(10,6))
plt.scatter(y_pred, residuals, color='blue', alpha=0.6)
plt.axhline(y=0, color='red', linestyle='--') # Line at 0 to show residuals centered
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()
```



### Prediction vs. Actual Plot

- This plot compares the predicted values with the actual values. A diagonal line represents perfect predictions (i.e., predicted = actual), and the points should ideally be close to this line.

|

```
plt.figure(figsize=(10,6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--') # Diagonal line for perfect predictions
plt.title('Prediction vs Actual')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.show()
```

